

# **Novel View Synthesis**

Joe Lin   Michael Song   Alexander Chien

February 23, 2024



# Intro + Problem Statement

Problem: I'm bored and I want to re-create 3D scenes from 2D images.

Solution: **Novel View Synthesis**, which aims to generate new scene perspectives given images at known viewpoints.

*Useful applications in:*

- Computational photography
- Virtual reality



# Classical Approach

**Light Fields** are 4D representation of radiance as a function of light ray position and direction.<sup>1</sup>

- **Light Slabs** ( $L(u, v, s, t)$ ) used to represent the light field.
- Less computational heavy and mappings from  $(x, y) \rightarrow (u, v, s, t)$  is a projection map.
- Create light field from views by mapping  $(x, y)$  to  $(u, v, s, t)$ .
- Views are essentially 2D slices of the 4D light field and to get novel views we can extract and resample slices of the light field (inverse mapping from  $(u, v, s, t)$  to  $(x, y)$ ).



**but wait...**

*I don't have enough photos from different perspectives 🤖. This classical approach isn't for me! 🤖 🚽*



# Classical Approach Challenges

- Sampling density must be high to avoid bluriness (i.e. many ground truth views are required).
- Heavy computation and memory usage.
- Observer is restricted to regions of space free of occluders.
- Illumination must be fixed.

Solution: Deep Learning approaches... **neural networks** learn mapping between sparse input views and high quality 3D scene representations.

# Neural Radiance Fields (NeRFs)<sup>2</sup>




# Algorithm Overview

Goal is to optimize a continuous **volumetric scene function**  $F_\theta$  where:

- input is a spatial location  $(x, y, z)$  and viewing direction  $(\theta, \varphi)$
- output is volume density  $\sigma$  and view-dependent emitted radiance

Basic NeRFs essentially encode 3D scene with a neural network.





# Analysis of Approach

Benefits:

- Capable of representing complex real-world geometry.
- Differentiable  $\Rightarrow$  can be optimized with gradient-based methods.





# Analysis of Approach

Benefits:

- Capable of representing complex real-world geometry.
- Differentiable  $\Rightarrow$  can be optimized with gradient-based methods.

**but wait...**



## Analysis of Approach

Benefits:

- Capable of representing complex real-world geometry.
- Differentiable  $\Rightarrow$  can be optimized with gradient-based methods.

**but wait...**

*I don't have enough computational power! I also want a good rendering  
in less than a day... 🤔*



## Analysis of Approach (cont.)

### Problems:

- Computationally expensive. Single scene can take days to train.
- Produces low FPS renders.
- Implicit representation (hard to interpret and edit scene).

### Solution: 3D Gaussian Splatting

# 3D Gaussian Splatting<sup>3</sup>



# Computer Graphics

**Rasterization** is the process of drawing graphical data onto the screen.

Objects are oftentimes represented by its polygonal faces (commonly triangles). Each of these polygons are decomposed into pixels and rasterized into a **raster image**.

*We can think of Gaussian Splatting as a rasterization technique.*




# Rasterizing Gaussians

$$G(x) = e^{(-\frac{1}{2})(x-\mu)^T \Sigma^{-1}(x-\mu)}$$

Essential parameters:

- position – mean,  $\mu$
- splat size – covariance matrix,  $\Sigma$
- color – spherical harmonics
- opacity –  $\alpha$

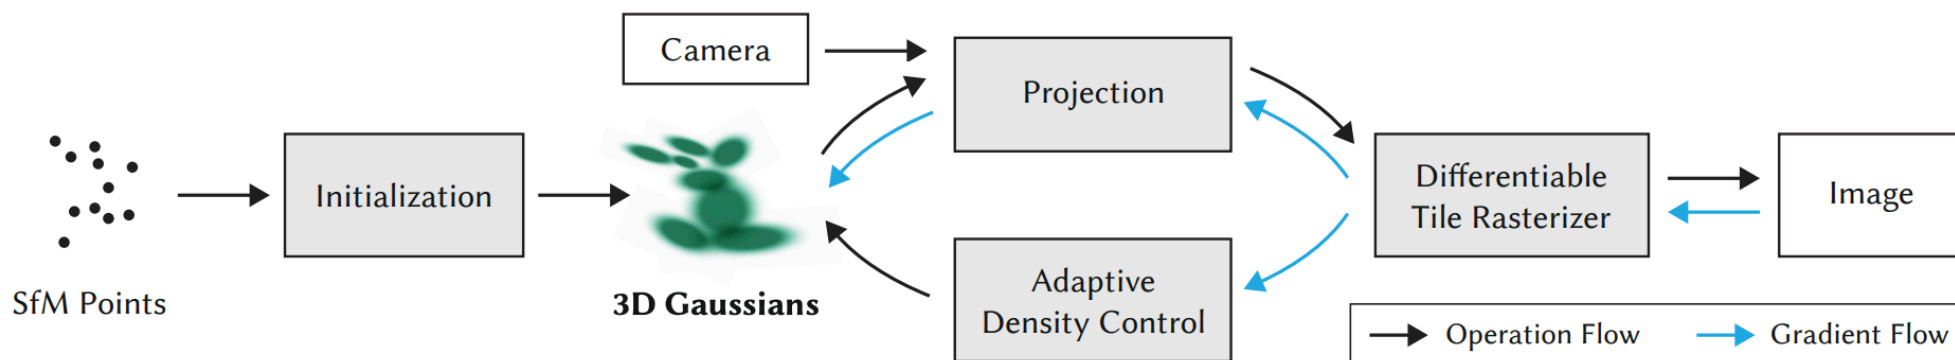




# Algorithm Overview

1. Given a set of images, estimate a pointcloud with **Structure from Motion (SfM)**.
2. Initialize isotropic Gaussians at each point.
3. Train 3D scene representation.
  - a. Rasterize Gaussians with a **differentiable rasterizer**.
  - b. Compute loss between raster image and ground truth image.
  - c. Adjust Gaussian parameters with SGD.
  - d. Densify or prune Gaussians.

# Algorithm Overview







# Differentiable Rasterizer

1. Project each Gaussian into 2D camera perspective.
  - a. Covariances in camera coordinates are  $\Sigma' = JW\Sigma W^T J^T$ 
    - i.  $J$  is the Jacobian of the affine approximation to the projective transform<sup>4</sup>
2. Sort Gaussians by order of depth.
3. Perform front-to-back blending of Gaussians at each pixel.



# Optimizing Gaussians

- Mean parameter  $\mu$
- Covariance parameter  $\Sigma$ 
  - *Initial Thought*: Optimize  $\Sigma$  directly.
    - **Problem: Doesn't guarantee  $\Sigma$  remains positive semi-definite.**
  - *What can we do instead?*
    - Gaussians can be seen as an **ellipsoid** rotated and stretched.
    - Decompose into  $\Sigma = RSS^T R^T$  and optimize  $R$  and  $S$ .<sup>5</sup>
- Spherical Harmonics coefficients  $c$
- Opacity parameter  $\alpha$



# Quaternions

Traditionally, rotations are represented with Euler angles (3 axis angles). In practice, rotation matrix  $R$  is encoded with **quaternions**.

$$\mathbb{H} = \{s + xi + yj + zk \mid s, x, y, z \in \mathbb{R}\}$$

where  $i, j, k$  obeys  $i^2 = j^2 = k^2 = ijk = -1, ij = k, ji = -k$

Quaternions prevent **gimbal lock**.



# Spherical Harmonics

Why not just RGB? Problem: No view dependence.

Spherical harmonics encode view-dependent colors as a linear combination of a basis of **harmonic functions**.<sup>5</sup>

$$\begin{aligned}
 & \begin{matrix} m = -2 & m = -1 & m = 0 & m = 1 & m = 2 \end{matrix} \\
 \ell = 0 & \quad c_0 \cdot \text{blue sphere} + \\
 \ell = 1 & \quad + c_2 \cdot \text{red sphere} + c_3 \cdot \text{yellow sphere} + c_4 \cdot \text{blue sphere} + \\
 \ell = 2 & \quad + c_5 \cdot \text{yellow sphere} + c_6 \cdot \text{blue sphere} + c_7 \cdot \text{blue sphere} + c_8 \cdot \text{yellow sphere} + c_9 \cdot \text{blue sphere} = \text{resulting sphere} \\
 & \quad \text{red} = \text{sigm}(\text{view sphere with } \theta, \phi) \cdot 255
 \end{aligned}$$



# Optimization Details

Use **Stochastic Gradient Descent** algorithm.

## Activation Functions

- Opacity parameters ( $\alpha$ ):  $\sigma(x) = \frac{1}{1+e^{-x}}$ 
  - smooth gradients
  - keep  $\alpha \in [0, 1)$
- Scale parameters ( $s$ ):  $\exp(x)$ 
  - similar reasons as above

**Initialization** of Gaussians: Estimate initial covariance matrix as an **isotropic** Gaussian with axes equal to the mean of the distances to the 3-NNs (recall the K-NN)

**Exponential decay scheduling** for positions.

## Loss Function

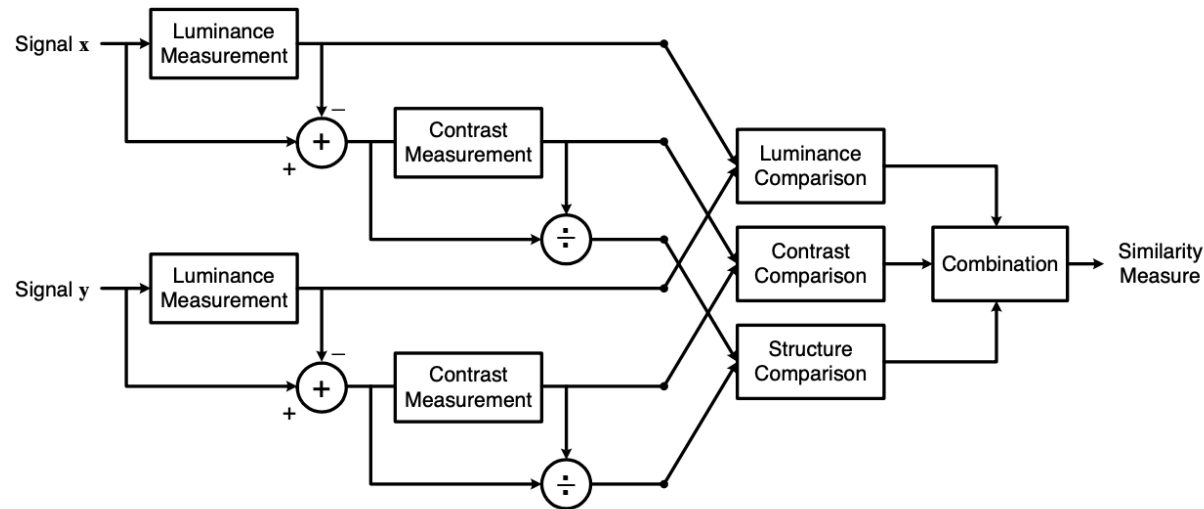
$$\mathcal{L} = (1 - \lambda)\mathcal{L}_1 + \lambda\mathcal{L}_{\text{D-SSIM}}$$

- Comparing predicted and ground truth intensities.
- $\lambda = 0.2$
- **L1 loss** ( $\mathcal{L}_1$ ): Mean Absolute Error
- D-SSIM ( $\mathcal{L}_{\text{D-SSIM}}$ ) = 1 - SSIM (**Structural Similarity Index Metric**)<sup>6</sup>
  - SSIM compares luminance, contrast, and structure



# Optimization Details (cont.)

$$\text{SSIM}(x, y) = \frac{(2\mu_x\mu_y + C_1)(2\sigma_{xy} + C_2)}{(\mu_x^2 + \mu_y^2 + C_1)(\sigma_x^2 + \sigma_y^2 + C_2)}$$





# Controlling Gaussians



Problem: Areas with missing geometric features (**under-reconstruction**; no Gaussians 🙄) and regions where Gaussians cover large areas in the scene (**over-reconstruction**): moral of the story is area not well-reconstructed.

- Both scenarios result in large view-space **positional gradients**.

## Solution: **Densification**

1. Small Gaussians in under-constructed regions: clone and move clone in direction of positional gradient. (Increase total volume of system and Gaussians)
2. Large Gaussians in regions of high variance: replace Gaussian with 2 new ones and divide scale by factor of  $\varphi = 1.6$  (value determined experimentally)
  - a. Initialize position by using original Gaussian as a probability density function for sampling.
  - b. Conserve total volume of system but increase number of Gaussians.

# Controlling Gaussians (cont.)







## Controlling Gaussians (cont.)

Another Issue Arises: Floaters close to input cameras cause unjustified increase in Gaussian density.

Solution: Set  $\alpha$  close to 0 every 3000 iterations.

Optimization increases  $\alpha$  for Gaussians where this is needed, while culling Gaussians with an alpha less than  $\epsilon_\alpha$ .

### Miscellaneous:

- Remove Gaussians large in world-space and those with a large footprint in view-space
- Gaussians remain primitives in Euclidean space
  - Don't require space compaction, warping, or projection strategies for distant and large Gaussians.





## Controlling Gaussians (cont.)

```
if IsRefinementIteration(i) then
  for all Gaussians  $(\mu, \Sigma, c, \alpha)$  in  $(M, S, C, A)$  do
    if  $\alpha < \epsilon$  or IsTooLarge( $\mu, \Sigma$ ) then           ▶ Pruning
      RemoveGaussian()
    end if
    if  $\nabla_p L > \tau_p$  then                             ▶ Densification
      if  $\|S\| > \tau_S$  then                             ▶ Over-reconstruction
        SplitGaussian( $\mu, \Sigma, c, \alpha$ )
      else                                             ▶ Under-reconstruction
        CloneGaussian( $\mu, \Sigma, c, \alpha$ )
      end if
    end if
  end for
end if
```




# Fast Differentiable Rasterizz

1. Split into  $16 \times 16$  tiles.
2. Cull 3D Gaussians with **frustum culling**.
3. Initialize Gaussians and give each a key.
4. Fast **radix sort** based on key (roughly  $O(n)$  algorithm).
5. Parallelize **alpha blending** per tile (one thread block per tile).
6. Stops thread block for tile when pixels reach certain alpha level.
7. Back-to-front traversal to compute gradients.



# Analysis of Approach

## Performance Metrics





## Analysis of Approach (cont.)

### Problems:

- Artifacts still exist (remember floaters?); especially elongated ones
  - Popping artifacts may arise due to trivial culling of rasterizz
- Significantly more memory than methods like NeRF

### Current/Future Works:

- Few shot 3D Gaussian Splatting
- Mesh reconstruction with Gaussians
- Reducing memory consumption



**BUT WAIT...**

*Ha. It's over. We don't have any more to offer... for the time being. 😏*



# Bibliography

1. Levoy M, Hanrahan P. Light Field Rendering. In: *Seminal Graphics Papers: Pushing the Boundaries, Volume 2*. 1st ed. Association for Computing Machinery; 2023. <https://doi.org/10.1145/3596711.3596759>
2. Mildenhall B, Srinivasan P P, Tancik M, Barron J T, Ramamoorthi R, Ng R. NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis. In: *ECCV*. ; 2020.
3. Kerbl B, Kopanas G, Leimkühler T, Drettakis G. 3D Gaussian Splatting for Real-Time Radiance Field Rendering. *ACM Transactions on Graphics*. 2023;42(4). <https://repo-sam.inria.fr/fungraph/3d-gaussian-splatting/>
4. Zwicker M, Pfister H, Baar J van, Gross M. EWA volume splatting. In: *Proceedings Visualization, 2001. VIS '01.* ; 2001:29-538. doi:10.1109/VISUAL.2001.964490

5. Yurkova K. A Comprehensive Overview of Gaussian Splatting. Published online December 2023. <https://towardsdatascience.com/a-comprehensive-overview-of-gaussian-splatting-e7d570081362>
6. Wang Z, Bovik A, Sheikh H, Simoncelli E. Image quality assessment: from error visibility to structural similarity. *IEEE Transactions on Image Processing*. 2004;13(4):600-612. doi:10.1109/TIP.2003.819861
7. Datta P. All about Structural Similarity Index (SSIM): Theory + Code in PyTorch. Published online September 2020. <https://medium.com/srm-mic/all-about-structural-similarity-index-ssim-theory-code-in-pytorch-6551b455541e>